

# Introduction to Neural Networks




SAPIENZA  
UNIVERSITÀ DI ROMA

Roberto Capobianco

Dipartimento di Ingegneria Informatica, Automatica e Gestionale


# Neural Networks: Motivation

- Initial goal: model biological neural systems
  - basic computational unit: neuron
  - ~86 billion neurons in the human nervous system
  - connected with  $\sim 10^{14k}$ - $10^{15}$  synapses
  - signals on axons interact multiplicatively with dendrites of other neurons based on some synaptic strength
- 




# Neural Networks: Implementation

- Diverged from biological model
  - engineered to achieve good results in ML tasks (different from real neurons!)
  - idea: synaptic strengths can be learned
  - model: dendrites carry signals that get summed in the cell body; if sum is above some threshold neuron fires
  - neurons fire with a frequency that depends on the activation function



# Artificial Neuron

- takes numerical **inputs** (x)
- has a **weight** associated to each input (w)
- has a **bias** in the form of an additional input 1 with weight b
- applies an activation function (f) to the weighted sum of inputs



$$\text{Output of neuron} = Y = f(w_1 \cdot X_1 + w_2 \cdot X_2 + b)$$

# Activation Function


- typically non-linear, aims at introducing non-linearity in the output of a neuron
- takes numbers as input
- performs a fixed mathematical operation on it

e.g., **softmax** function: takes a vector of arbitrary real-valued scores (in  $z$ ) and squashes it to a vector of values between zero and one that sum to one.

# Types of Activation Function


- sigmoid (bad!): takes a real-valued input and squashes it to range between 0 and 1

$$\sigma(x) = 1 / (1 + \exp(-x))$$




- 
- tanh: takes a real-valued input and squashes it to the range  $[-1, 1]$

$$\tanh(x) = 2\sigma(2x) - 1$$





- ReLU: ReLU stands for Rectified Linear Unit. It takes a real-valued input and thresholds it at zero (replaces negative values with zero)

$$f(x) = \max(0, x)$$





# Role of Bias



# Neural Network Architecture

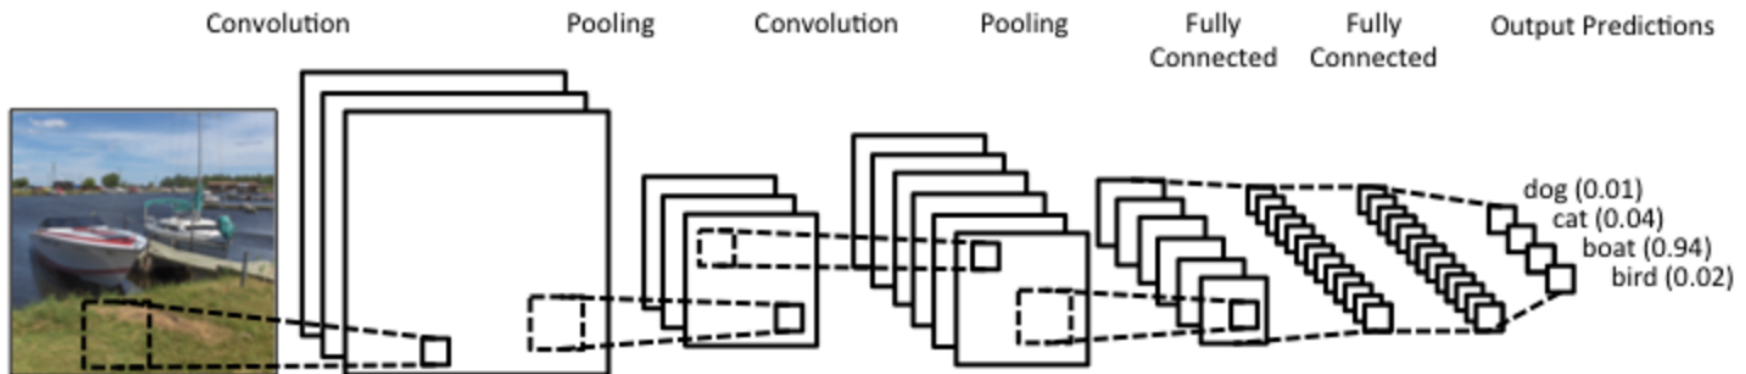
- regular neural networks are neurons connected in an acyclic graph
- 1 or more layers
- typically fully connected layers (no connection inside the same layer)
- output layer typically without activation function
- naming convention: input layer is not counted
  - single-layer networks directly map input to output





# Alternative Types of Neural Networks

- **Convolutional Neural Networks**
- unit connectivity pattern inspired by the organization of the visual cortex
- units respond to stimuli in a restricted region of space known as the receptive field
- receptive fields partially overlap, over-covering the entire visual field



# Alternative Types of Neural Networks

- **Recurrent Neural Networks**
- connections between unit form a directed cycle
- propagate data forward AND backward
- use memory to process sequence of inputs
- useful for tasks like speech recognition and sequence processing

# Training Neural Networks

- **Backpropagation**

- 1) Initialize all edge weights randomly
- 2) For every input, until error < threshold:
  - a) provide it to the neural network and observe output
  - b) compare output with desired output / label
  - c) propagate the error back to the previous layer


## Additional resources:

<https://github.com/rasbt/python-machine-learning-book/blob/master/faq/visual-backpropagation.md>


<https://ujjwalkarn.me/2016/08/09/quick-intro-neural-networks/>

3D visualization of neural networks: <http://scs.ryerson.ca/~aharley/vis/fc/>

# Backpropagation




Backpropagation  
+  
Weights Adjusted




# Backpropagation

Correct Output



# Convolutional Neural Network



# Convolutional Layer

- 'eyes' of the neural network
- neurons look for specific features
- to calculate convolution at a particular location (x, y):
  - extract a chunk from the image centered at location (x,y)
  - multiply the values in this chunk element-by-element with the convolution filter (sized as the chunk)
  - add them all to obtain a single output
- weights and biases of neurons corresponding to one filter are shared

7	2	3	3	8
4	5	3	8	4
3	3	2	8	4
2	8	7	2	7
5	4	4	5	4


\*

1	0	-1
1	0	-1
1	0	-1

=


6		

$7 \times 1 + 4 \times 1 + 3 \times 1 + 2 \times 0 + 5 \times 0 + 3 \times 0 + 3 \times -1 + 3 \times -1 + 2 \times -1 = 6$



# Stride, Multiple Filters and Padding



- stride: how many pixels we use to slide the filter
- after each convolution the output reduces in size (filter 3x3x3 → activation map 30x30x32)
- pad zeros to the boundary of the input layer such that the output is the same size as input layer













# Hierarchical Features & Max Pooling

- pooling reduces spatial size and number of parameters
- avoids overfitting
- commonly: max pooling



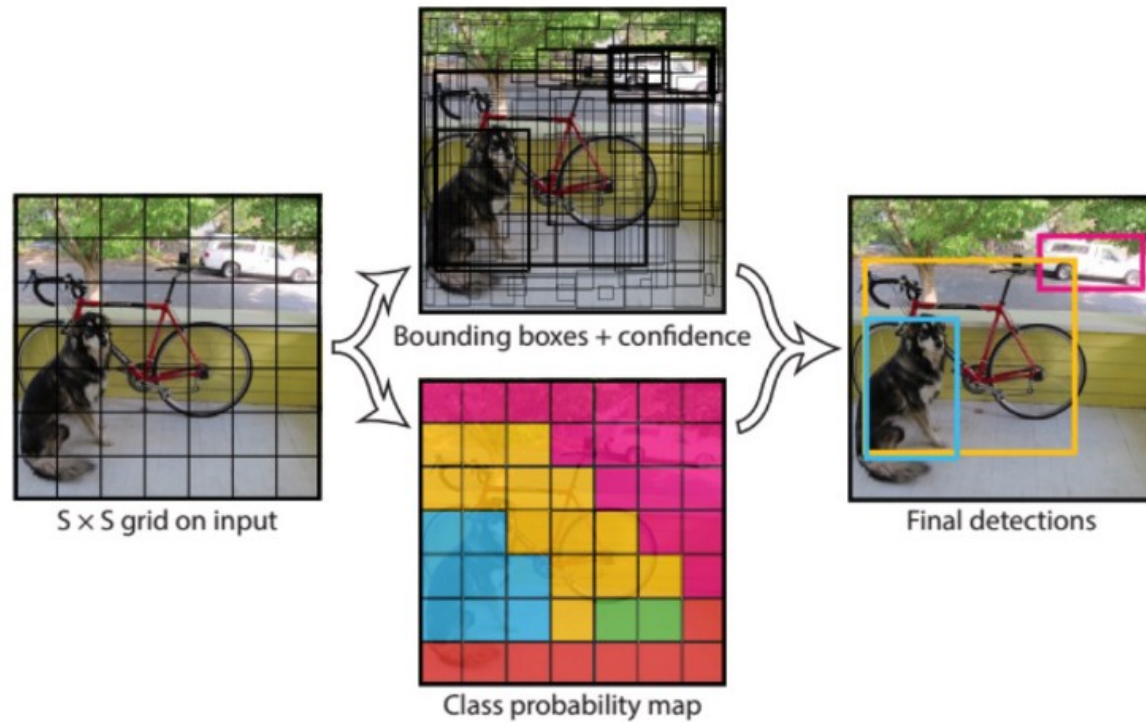
# Data Augmentation



# YOLO


- YOLO: You Only Look Once
- state-of-the-art real-time object detection system
- <https://pjreddie.com/darknet/yolo/>
-

# YOLO: Idea



# YOLO: Architecture & Limitations

- YOLO: You Only Look Once
- state-of-the-art real-time object detection system
- <https://pjreddie.com/darknet/yolo/>
- 



**Figure 3: The Architecture.** Our detection network has 24 convolutional layers followed by 2 fully connected layers. Alternating  $1 \times 1$  convolutional layers reduce the features space from preceding layers. We pretrain the convolutional layers on the ImageNet classification task at half the resolution ( $224 \times 224$  input image) and then double the resolution for detection.

# Homework

Download and install MxNet or TensorFlow (choose if you want to use CPU or also GPU)

<https://mxnet.incubator.apache.org/install/index.html>

[https://www.tensorflow.org/install/install\\_linux](https://www.tensorflow.org/install/install_linux)

Follow the MNIST tutorial:

MxNet

<https://mxnet.incubator.apache.org/tutorials/python/mnist.html>

TensorFlow

[https://www.tensorflow.org/versions/r1.1/get\\_started/mnist/beginners](https://www.tensorflow.org/versions/r1.1/get_started/mnist/beginners)

[https://www.tensorflow.org/versions/r1.2/get\\_started/mnist/pros](https://www.tensorflow.org/versions/r1.2/get_started/mnist/pros)

Change parameters, layers, activation function, etc. of your network and compare results.

Write a short summary / report to discuss at correction time.